

MAST90104: A First Course in Statistical Learning

Week 4 Workshop and Lab **Solutions**

Workshop questions

1. Suppose that X is a random variable with density function, f , given by

$$f(x) = \sum_{i=0}^{\infty} p(i)g(x; i)$$

where $p(0), p(1), \dots$ is a discrete probability mass function on $\{0, 1, \dots\}$ and each $g(x; i)$ is a probability density function. Suppose that $\mu(i), \sigma^2(i), M(t; i)$ are the mean, variance and moment generating function for the density $g(x; i)$. Let $M(t)$ be the moment generating function of X . Suppose also that N is a random variable with probability mass function $p(i), i = 0, 1, \dots$. Show that

(a) $E(X) = E(\mu(N))$

Solution:

$$\begin{aligned} E(X) &= \int_{-\infty}^{\infty} xf(x) dx \\ &= \int_{-\infty}^{\infty} x \sum_{i=0}^{\infty} p(i)g(x; i) dx \\ &= \sum_{i=0}^{\infty} \int_{-\infty}^{\infty} xg(x; i) dx p(i) \\ &= \sum_{i=0}^{\infty} \mu(i)p(i) \\ &= E(\mu(N)) \end{aligned}$$

(b) $\text{var}(X) = E(\sigma^2(N)) + \text{var}(\mu(N))$

Solution:

$$\begin{aligned} \text{var}(X) &= E(X^2) - (E(X))^2 \\ &= \int_{-\infty}^{\infty} x^2 \sum_{i=0}^{\infty} p(i)g(x; i) dx - (E(X))^2 \\ &= \sum_{i=0}^{\infty} \int_{-\infty}^{\infty} x^2 g(x; i) dx p(i) - (E(X))^2 \\ &= \sum_{i=0}^{\infty} (\sigma^2(i) + \mu^2(i))p(i) - (E(X))^2 \\ &= E(\sigma^2(N)) + E(\mu^2(N)) - (E(\mu(N)))^2 \\ &= E(\sigma^2(N)) + \text{var}(\mu(N)) \end{aligned}$$

(c) $M(t) = E(M(t; N))$.

Solution:

$$\begin{aligned}
M(t) &= E(e^{tX}) \\
&= \int_{-\infty}^{\infty} e^{tx} \sum_{i=0}^{\infty} p(i)g(x; i) dx \\
&= \sum_{i=0}^{\infty} \int_{-\infty}^{\infty} e^{tx} g(x; i) dx p(i) \\
&= \sum_{i=0}^{\infty} M(t; i)p(i) \\
&= E(M(t; N))
\end{aligned}$$

(Hint: You may assume that interchange of infinite sums and integrals is justified.)

2. Let y_1, \dots, y_n be an i.i.d. normal sample. Show that

$$\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i \text{ and } s^2 = \frac{1}{n-1} \sum_{i=1}^n (y_i - \bar{y})^2$$

are independent. (*Hint:* Express them as a random “vector” and quadratic form respectively.)

Solution: Let $\mathbf{1}$ be the vector made up entirely of 1's, then

$$\begin{aligned}
\bar{y} &= \frac{1}{n} \mathbf{1}^T \mathbf{y} \\
\mathbf{y} - \bar{y} \mathbf{1} &= (I - \frac{1}{n} \mathbf{1} \mathbf{1}^T) \mathbf{y} \\
s^2 &= \frac{1}{n-1} [(I - \frac{1}{n} \mathbf{1} \mathbf{1}^T) \mathbf{y}]^T (I - \frac{1}{n} \mathbf{1} \mathbf{1}^T) \mathbf{y} \\
&= \frac{1}{n-1} \mathbf{y}^T (I - \frac{1}{n} \mathbf{1} \mathbf{1}^T)^T (I - \frac{1}{n} \mathbf{1} \mathbf{1}^T) \mathbf{y} \\
&= \frac{1}{n-1} \mathbf{y}^T (I - \frac{1}{n} \mathbf{1} \mathbf{1}^T) \mathbf{y}
\end{aligned}$$

noting that $I - \frac{1}{n} \mathbf{1} \mathbf{1}^T$ is symmetric and idempotent. It is now easy to check that $B = \frac{1}{n} \mathbf{1}^T$ and $A = \frac{1}{n-1} (I - \frac{1}{n} \mathbf{1} \mathbf{1}^T)$ satisfy $BVA = 0$, where $V = \sigma^2 I = \text{Var } \mathbf{y}$, whence $\bar{y} = B\mathbf{y}$ and $s^2 = \mathbf{y}^T A \mathbf{y}$ are independent.

The alternative approach is to notice that \bar{y} and s^2 are just the usual estimates of β and σ^2 for the linear model $\mathbf{y} = \mathbf{1}\beta + \varepsilon$, with $\beta = \mu = E(y_i)$ and $\text{Var } \varepsilon = \sigma^2 I$. (This is sometimes called the *null* model.)

3. An online survey collects data on factors that affect a person's pay rate (per hour). The table below shows pay rate (**pay**) and number of years of education (**yrEdu**) of five participants.

id	pay	yrEdu
1	7.06	9
2	18.93	12
3	20.17	12
4	29.58	16
5	33.90	20

- (a) Let x_i and y_i denote the years of education and pay rate of individual i . We want to fit the model $y_i = \beta_0 + \beta_1 x_i + \epsilon_i$. Given that $\sum_i x_i^2 = 1025$, $\bar{x} = 13.8$, $\bar{y} = 21.928$, $\sum_i x_i y_i = 1684.02$, find the least squares estimates of β_0, β_1

Solution This is a simple linear regression model, so using the formulas in the lecture notes:

$$\begin{aligned}
\hat{\beta}_1 &= \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} \\
&= \frac{\sum_{i=1}^n x_i y_i - n \bar{x} \bar{y}}{\sum_{i=1}^n x_i^2 - n \bar{x}^2} \\
&= \frac{1684.02 - 5 \times 13.8 \times 21.928}{1025 - 5 \times 13.8^2} \\
&= 2.348736
\end{aligned}$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x} = -10.48456$$

- (b) Suppose we have calculate the least squares estimators $\hat{\beta}_0$ and $\hat{\beta}_1$ in R. Consider the following R commands and output

```
> pay = c(7.06,18.93,20.17,29.58,33.9)
> yrEdu = c(9,12,12,16,20)
> e = pay - (betahat0 + betahat1*yrEdu)
> t(e)%*%e
[,1]
[1,] 33.41216
```

Calculate the sample variance s^2 .

Solution $s^2 = SS_{\text{Res}}/(n - p) = 33.41216/(5 - 2) = 11.13739$

- (c) Estimate the pay rate of a person with 14 years of education.

Solution $\mathbf{t} = (1, 14)^T$. Then estimated pay rate is $\mathbf{t}^T \hat{\boldsymbol{\beta}} = 22.39775$.

- (d) The leverage of the data points are given as

```
> model1_leverage
[1] 0.5164835 0.2445055 0.2445055 0.2664835 0.7280220
```

Calculate the standardised residual for the 3rd observation.

Solution $e_3 = 20.17 - \hat{\beta}_0 - \hat{\beta}_1 \times 12 = 2.46973$.

$z_3 = e_3 / \sqrt{s^2(1 - H_{33})} = 2.46973 / \sqrt{11.13739 \times (1 - 0.2445055)} = 0.8514166$.

Practical questions

Attempt the exercises below.

1. Last week you wrote a program to calculate $h(x, n)$, the sum of a finite geometric series. Turn this program into a *function* that takes two arguments, x and n , and returns $h(x, n)$.

Make sure you deal with the case $x = 1$.

Solution:

```
series_sum <- function(x, n) {  
  # sum of x^k for k = 0, ..., n  
  if (x == 1) {  
    return(n + 1)  
  } else {  
    return((x^(n+1) - 1)/(x - 1))  
  }  
}
```

2. Consider the following program

```
# clear the workspace  
rm(list=ls())  
  
random.sum <- function(n) {  
  # sum of n random numbers  
  x[1:n] <- ceiling(10*runif(n))  
  cat("x:", x[1:n], "\n")  
  return(sum(x))  
}
```

Below are the output of the function for $n = 10$ and $n = 5$

```
> x <- rep(100, 10)  
> show(random.sum(10))  
x: 6 10 7 5 8 6 5 10 9 4  
[1] 70  
> show(random.sum(5))  
x: 8 9 4 5 10  
[1] 536
```

Explain what is going wrong and how you would fix it.

Solution

The problem is is the line

```
x[1:n] <- ceiling(10*runif(n))
```

One way to fix it is

```
random.sum.fix <- function(n) {  
  # sum of n random numbers  
  x <- ceiling(10*runif(n))  
  cat("x:", x[1:n], "\n")  
  return(sum(x))  
}
```

3. In this question we simulate the rolling of a die. To do this we use the function `runif(1)`, which returns a 'random' number in the range (0,1). To get a random integer in the range $\{1, 2, 3, 4, 5, 6\}$, we use `ceiling(6*runif(1))`, or if you prefer, `sample(1:6,size=1)` will do the same job.

- (a) Suppose that you are playing the gambling game of the Chevalier de Méré. That is, you are betting that you get at least one six in four throws of a die. Write a program that simulates one round of this game and prints out whether you win or lose.

Check that your program can produce a different result each time you run it.

Solution:

```
win <- FALSE
for (i in 1:4) {
  if (sample(1:6, size = 1) == 6) {
    win <- TRUE
  }
}
if (win) {
  print("win")
} else {
  print("lose")
}

## [1] "win"
```

- (b) Turn the program that you wrote in part (a) into a function `sixes`, which returns `TRUE` if you obtain at least one six in n rolls of a fair die, and returns `FALSE` otherwise. That is, the argument is the number of rolls n , and the value returned is `TRUE` if you get at least one six and `FALSE` otherwise.

How would you give n the default value of 4?

Solution:

```
sixes <- function(n = 4) {
  # plays the game of the Chevalier de Mere
  # returns TRUE if at least one six in n rolls
  # returns FALSE otherwise
  win <- FALSE
  for (i in 1:n) {
    if (sample(1:6, size = 1) == 6) {
      return(TRUE)
    }
  }
  return(FALSE)
}
```

Here is a vectorised version.

```
sixes <- function(n = 4) {
  sum(sample(1:6, size = n, replace = TRUE) == 6) > 0
}
```

- (c) Now write a program that uses your function `sixes` from part (b), to simulate N plays of the game (each time you bet that you get at least one six in n rolls of a fair die). Your program should then determine the proportion of times you win the bet. This proportion is an estimate of the *probability* of getting at least one six in n rolls of a fair die.

Run the program for $n = 4$ and $N = 100, 1000$, and 10000 , conducting several runs for each N value. How does the *variability* of your results depend on N ?

The probability of getting no 6's in n rolls of a fair die is $(5/6)^n$, so the probability of getting at least one is $1 - (5/6)^n$. Modify your program so that it calculates the theoretical probability as well as the simulation estimate and prints the difference between them. How does the *accuracy* of your results depend on N ?

Solution:

```

p_estimate <- function(N, n = 4) {
  # proportion of wins in N runs of sixes(n)
  total_wins <- 0
  for (i in 1:N) {
    if (sixes(n)) total_wins <- total_wins + 1
  }
  return(total_wins/N)
}
p_accuracy <- function(N, n = 4) {
  # accuracy of p_estimate
  total_wins <- 0
  for (i in 1:N) {
    if (sixes(n)) total_wins <- total_wins + 1
  }
  return(total_wins/N - 1 + (5/6)^n)
}

```

- (d) In part (c), instead of processing the simulated runs as we go, suppose we first store the results of every game in a file, then later postprocess the results.

Write a program to write the result of all N runs to a textfile `sixes_sim.txt`, with the result of each run on a separate line. For example, the first few lines of the textfile could look like

```

# TRUE
# FALSE
# FALSE
# TRUE
# FALSE
# .
# .

```

Now write another program to read the textfile `sixes_sim.txt` and again determine the proportion of bets won.

This method of saving simulation results to a file is particularly important when each simulation takes a very long time (hours or days), in which case it is good to have a record of your results in case of a system crash.

Solution:

```

sixes_sim <- function(N, n = 4) {
  # runs sixes(n) N times and saves the results in "sixes_sim.txt"
  cat(file="sixes_sim.txt") # deletes contents of file
  for (i in 1:N) {
    cat(file = "sixes_sim.txt", sixes(n), "\n", append = TRUE)
  }
}
sixes_sim(100)
results <- scan("sixes_sim.txt", what = TRUE)
mean(results)

## [1] 0.59

```

4. Let $\mathbf{y} = \begin{bmatrix} y_1 & y_2 \end{bmatrix}^T$ be a normal random vector with mean and variance

$$\boldsymbol{\mu} = \begin{bmatrix} 2 \\ 4 \end{bmatrix}, \quad V = \begin{bmatrix} 2 & 0 \\ 0 & 2 \end{bmatrix}.$$

Let

$$A = \frac{1}{4} \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}.$$

From Theorem 3.9 we know that $\mathbf{y}^T \mathbf{A} \mathbf{y}$ follows a χ^2 distribution with degree of freedom 1 and noncentrality parameter $\lambda = 4.5$.

Solution: We first need to verify whether $\mathbf{A}\mathbf{V}$ is idempotent

```
> mu = c(2,4);
> V = 2*diag(2)
> A = 1/4*matrix(c(1,1,1,1),2,2);
> A
      [,1] [,2]
[1,] 0.25 0.25
[2,] 0.25 0.25
>
> (AV = A%*%V)
      [,1] [,2]
[1,] 0.5 0.5
[2,] 0.5 0.5
> AV%*%AV
      [,1] [,2]
[1,] 0.5 0.5
[2,] 0.5 0.5
>
> 1/2*mu%*%A%*%mu
      [,1]
[1,] 4.5
```

(a) Generate $n = 1000$ samples $\{\mathbf{y}^{(1)}, \dots, \mathbf{y}^{(n)}\}$ from $MVN(\boldsymbol{\mu}, V)$.

```
> n = 1000
> ysample = mvrnorm(n,mu,V)
> dim(ysample)
[1] 1000 2
```

(b) Compute $\mathbf{y}^T \mathbf{A} \mathbf{y}$ for all $\mathbf{y}^{(i)}$ that we generated.

```
> quadform = function(y, A) t(y) %*% A %*% y
> quadsampleA = apply(ysample,1,quadform, A= A)
> str(quadsampleA)
 num [1:1000] 5.4 15.87 4.77 8.24 12.51 ...
```

(c) Plot the histogram of the $\mathbf{y}^T \mathbf{A} \mathbf{y}$ values that we have computed.

```
> hist(quadsampleA, col = 'coral1')
```

(d) Now generate n samples from $\chi^2_{1,4.5}$ distribution using `rchisq()`.

```
> chiA = rchisq(n,1,mu%*%A%*%mu)
```

(e) Plot the histogram of the generated samples on the same graph with the histogram in part (c).

The two histograms should overlap.

```
> hist(chiA,add = T, col = 'lightcyan')
> # not very nice so improve it a little bit by making the color transparent
> # first get the red, green and blue values needed for the rgb() command
> col2rgb(c("coral1","lightcyan"))
      [,1] [,2]
red      255  224
green    114  255
blue      86  255
> hist(quadsampleA, col = rgb(255,114,86,max = 255,alpha = 125))
> # because rgb color are defined in range 0-255
> hist(chiA, col = rgb(225,255,255,max = 255,alpha = 125),add = T)
```