

# MAST90104: A First Course in Statistical Learning

## Week 7 Practical and Workshop - Solution

### 1 Practical questions

1. The data *teengamb* from the **faraway** package contains data from a survey in Britain to study teenage gambling. The variables are

- **sex**: 0=male, 1=female
- **status**: Socioeconomic status score based on parents' occupation income in pounds per week
- **verbal**: verbal score in words out of 12 correctly defined
- **gamble**: expenditure on gambling in pounds per year

We can import the data by

```
data(teengamb)
```

We are interested in predicting **gamble** using the other variables. Implement the following variable selection methods to determine the “best” model.

(Note that here we are ignoring the fact gender can be treated as factor)

- (a) Backward elimination

#### Solution

```
> install.packages('faraway')
> library(faraway)
> data(teengamb)
> nullmodel <- lm(gamble ~ 1 ,data = teengamb)
> fullmodel <- lm(gamble~., data = teengamb)
> # backward selection
> drop1(fullmodel,scope= ~ ., test="F")
Single term deletions
```

Model:

		Df	Sum of Sq	RSS	AIC	F value	Pr(>F)
			21624	298.18			
<none>							
sex	1	3735.8	25360	303.67	7.2561	0.01011	*
status	1	17.8	21642	296.21	0.0345	0.85349	
income	1	12056.2	33680	317.00	23.4169	1.792e-05	***
verbal	1	955.7	22580	298.21	1.8563	0.18031	

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
> model2 <- lm(gamble ~ sex + income + verbal,data = teengamb)
> drop1(model2, scope = ~.,test = 'F')
Single term deletions
```

Model:

		Df	Sum of Sq	RSS	AIC	F value	Pr(>F)
			21642	296.21			
<none>							
sex	1	5787.9	27429	305.35	11.5001	0.001502	**
income	1	13236.1	34878	316.64	26.2990	6.644e-06	***
verbal	1	1139.8	22781	296.63	2.2646	0.139667	

```

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
> model3 <- lm(gamble ~ sex + income,data = teengamb)
> drop1(model3, scope = ~., test = 'F')
Single term deletions

Model:
gamble ~ sex + income
      Df Sum of Sq  RSS    AIC F value    Pr(>F)
<none>            22781 296.63
sex      1      5227.3 28009 304.34  10.096 0.002717 **
income   1     15309.8 38091 318.79  29.569 2.245e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
> # stop

```

(b) Forward selection

**Solution:**

```

> # forward selection
> add1(nullmodel,scope = ~. +sex + status + income + verbal,test= 'F')
Single term additions

Model:
gamble ~ 1
      Df Sum of Sq  RSS    AIC F value    Pr(>F)
<none>            45689 325.34
sex      1      7598.4 38091 318.79  8.9766 0.004437 **
status   1       116.2 45573 327.22  0.1147 0.736438
income   1     17680.9 28009 304.34 28.4070 3.045e-06 ***
verbal    1      2212.5 43477 325.00  2.2900 0.137202
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
> model2 <- lm(gamble~income,data = teengamb)
> add1(model2, scope = ~. +sex + status + verbal,test= 'F')
Single term additions

```

```

Model:
gamble ~ income
      Df Sum of Sq  RSS    AIC F value    Pr(>F)
<none>            28009 304.34
sex      1      5227.3 22781 296.63 10.0960 0.002717 **
status   1       719.8 27289 305.11  1.1605 0.287228
verbal    1      579.1 27429 305.35  0.9290 0.340385
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
> model3 <- lm(gamble ~ sex + income, data = teengamb)
> add1(model3, scope = ~. + status + verbal,test= 'F')
Single term additions

```

```

Model:
gamble ~ sex + income
      Df Sum of Sq  RSS    AIC F value Pr(>F)
<none>            22781 296.63
status   1       201.82 22580 298.21  0.3843 0.5386
verbal    1     1139.78 21642 296.21  2.2646 0.1397

```

(c) Stepwise selection

**Solution**

```
> # stepwise selection
> step(nullmodel, scope = ~. + sex + status + income + verbal)
Start:  AIC=325.34
gamble ~ 1
```

	Df	Sum of Sq	RSS	AIC
+ income	1	17680.9	28009	304.34
+ sex	1	7598.4	38091	318.79
+ verbal	1	2212.5	43477	325.00
<none>			45689	325.34
+ status	1	116.2	45573	327.22

```
Step:  AIC=304.34
gamble ~ income
```

	Df	Sum of Sq	RSS	AIC
+ sex	1	5227.3	22781	296.63
<none>			28009	304.34
+ status	1	719.8	27289	305.11
+ verbal	1	579.1	27429	305.35
- income	1	17680.9	45689	325.34

```
Step:  AIC=296.63
gamble ~ income + sex
```

	Df	Sum of Sq	RSS	AIC
+ verbal	1	1139.8	21642	296.21
<none>			22781	296.63
+ status	1	201.8	22580	298.21
- sex	1	5227.3	28009	304.34
- income	1	15309.8	38091	318.79

```
Step:  AIC=296.21
gamble ~ income + sex + verbal
```

	Df	Sum of Sq	RSS	AIC
<none>			21642	296.21
- verbal	1	1139.8	22781	296.63
+ status	1	17.8	21624	298.18
- sex	1	5787.9	27429	305.35
- income	1	13236.1	34878	316.64

```
Call:
lm(formula = gamble ~ income + sex + verbal, data = teengamb)
```

```
Coefficients:
(Intercept)      income          sex        verbal
    24.139         4.898       -22.960        -2.747
```

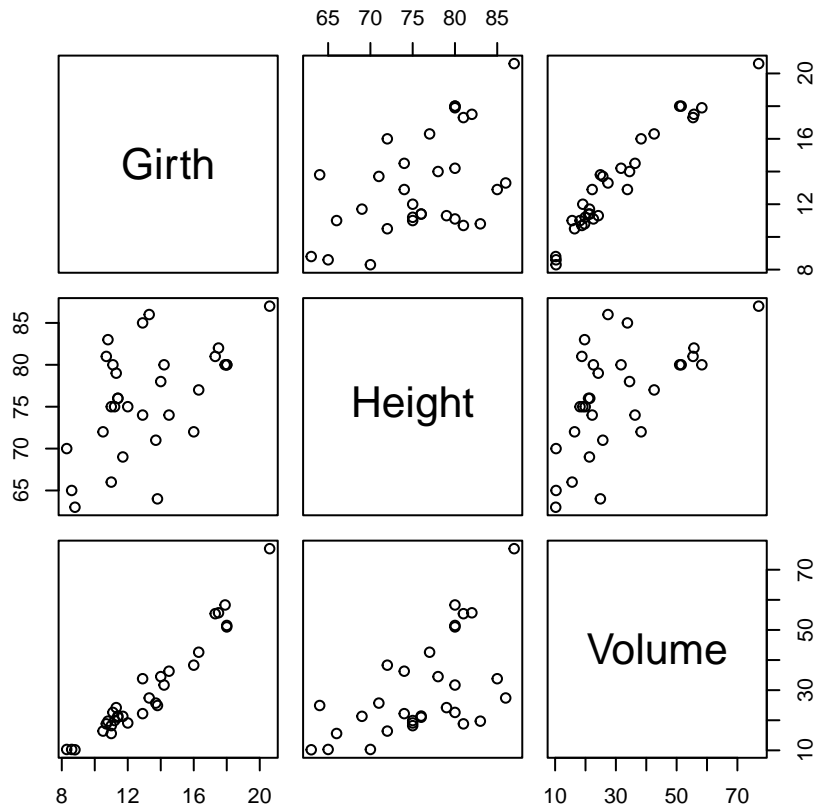
Comment on the AIC of the final model chosen by the 3 methods.

**Solution:** Backward and forward selection chose the same variables (**sex** and **income**) while stepwise selection final model has one more parameters, **verbal**. The AIC of the 2 models are very close to each other (296.63 vs 296.21).

2. Load and examine the dataset **trees** using

```
data(trees)
?trees
```

```
pairs(trees)
```



We will model the volume of a black cherry tree as a function of its girth and height.

- (a) By calculating  $R(\gamma_1|\gamma_2)$  and  $SS_{Res}$  from the data  $\mathbf{y}$  and design matrix  $X$ , use an F test to determine if including the variable **Height** significantly improves the model fitted using only **Girth** (and an intercept).

**Solution**

```
> data(trees)
> # a.
> y <- trees$Volume
> n <- length(y)
> # model with both Girth and Height
> X <- cbind(1, trees$Girth, trees$Height)
> betahat <- solve(t(X) %*% X, t(X) %*% y)
> SS_res <- sum((y - X %*% betahat)^2)
[1] 421.9214
> SS_reg <- sum((X %*% betahat)^2)
> # model with only Girth + intercept
> X2 <- X[, -3]
> betahat2 <- solve(t(X2) %*% X2, t(X2) %*% y)
> SS_reg2 <- sum((X2 %*% betahat2)^2)
> R_g1g2 <- SS_reg - SS_reg2
[1] 102.3812
> Fstat <- (R_g1g2/1)/(SS_res/(n - 3))
[1] 6.79433
> pf(Fstat, 1, n - 3, lower.tail = F)
[1] 0.01449097
```

Repeat the test using the `lm` and `anova` commands, to see if you get the same numbers.

```
> model1 <- lm(Volume ~ Girth, data = trees)
> model2 <- lm(Volume ~ Girth + Height, data = trees)
> anova(model1, model2)
Analysis of Variance Table

Model 1: Volume ~ Girth
Model 2: Volume ~ Girth + Height
    Res.Df    RSS Df Sum of Sq    F Pr(>F)
1       29 524.30
2       28 421.92  1    102.38 6.7943 0.01449 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- (b) Add variables `Girth` squared and `Girth` squared times `Height` to the model, then use stepwise selection to simplify the model. (You can use `step` for this step.). Comment on the form of your final model.

### Solution

```
> trees$GirthSq <- trees$Girth^2
> model <- lm(Volume ~ Girth + Height + GirthSq + GirthSq*Height, data = trees)
> model <- step(model, scope = ~ .)
Start: AIC=64.36
Volume ~ Girth + Height + GirthSq + GirthSq * Height
```

	Df	Sum of Sq	RSS	AIC
- Girth	1	0.2288	179.27	62.402
- Height:GirthSq	1	6.9694	186.01	63.547
<none>			179.04	64.363

```
Step: AIC=62.4
Volume ~ Height + GirthSq + Height:GirthSq
```

	Df	Sum of Sq	RSS	AIC
<none>			179.27	62.402
+ Girth	1	0.229	179.04	64.363
- Height:GirthSq	1	40.164	219.44	66.669

Note that R will not attempt to drop `GirthSq` and `Height` if the interaction term `GirthSq*Height` is still in the model.

- (c) Use diagnostic plots to check the fit of your final model.
- (d) What transformation might be indicated from the plot of residuals versus fitted values? Transform all variables with this transformation. What might the appropriate model be? Fit it and comment on the resulting residuals.

### Solution

The plot of residuals versus fitted values show a fan-shape pattern, signalling that the errors do not have constant variance. We may want to transform the response variable. The Scale-Location plot show some increase in  $\sqrt{|\text{Standardized residuals}|}$  as fitted values increases, so a log transformation may be appropriate.

We consider 2 models: one with `Height` and `Girth`, and one where we also apply a log transformation to `Height` and `Girth`. The diagnostics plots of the 2 models (Figure 2 and Figure 3 ) suggest the log-log model is better. The residuals plot in `model3` looks slightly better than the initial model, however there is still some heteroskedasticity shown in the plot.

```
> model2 <- lm(log(Volume) ~ Height + Girth, data = trees)
> model3 <- lm(log(Volume) ~ log(Height) + log(Girth), data = trees)
```

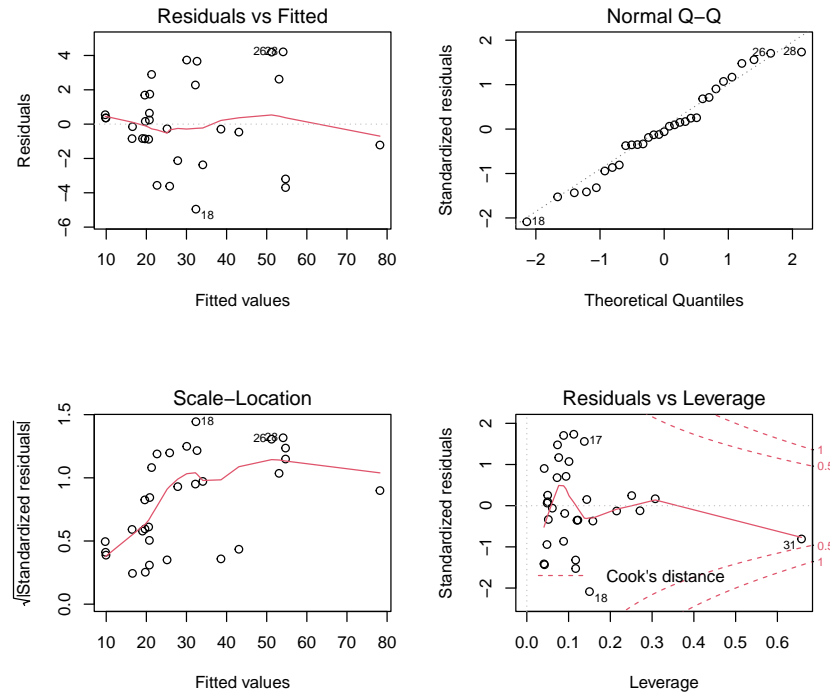


Figure 1: Diagnostic plots of the final model

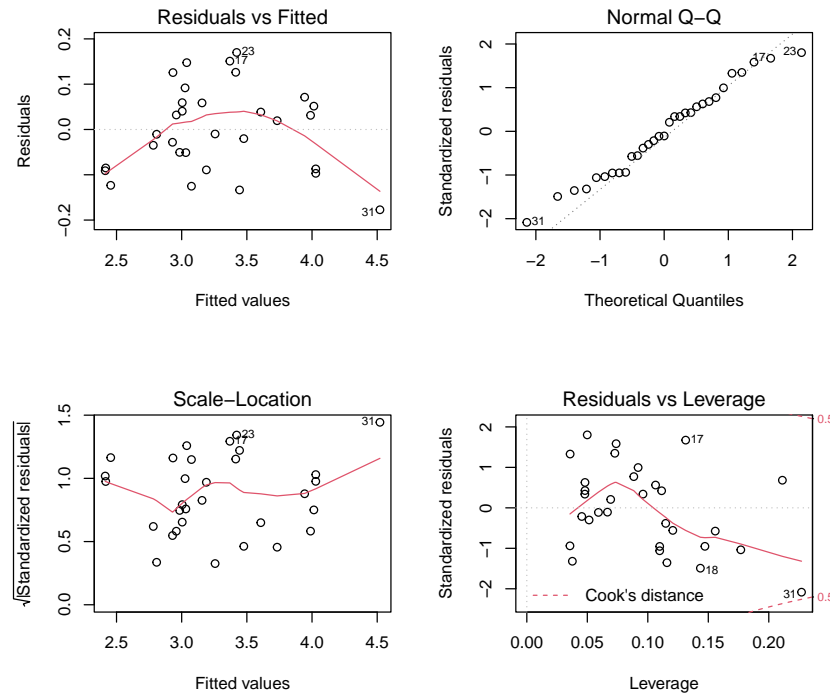


Figure 2: Diagnostic plots of model2

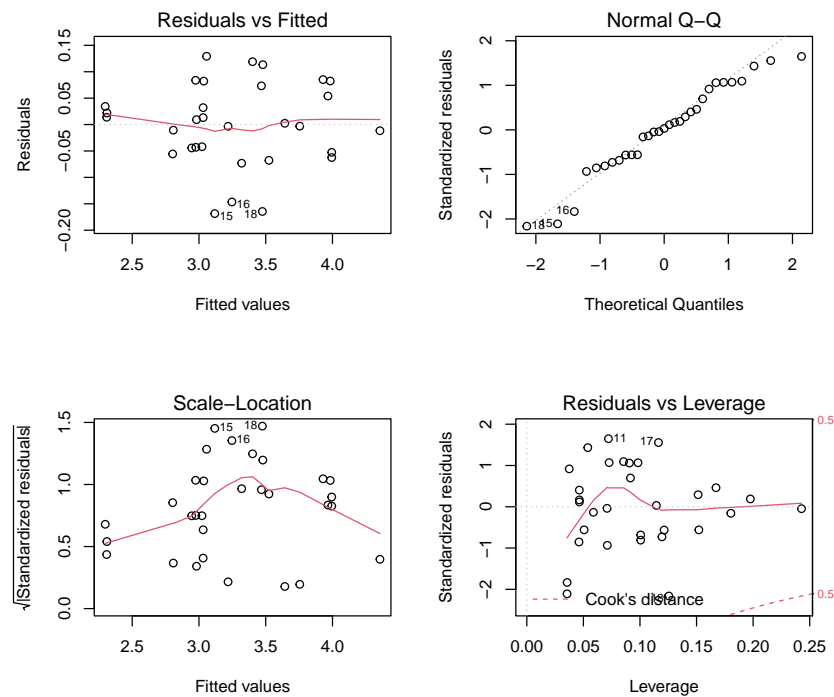


Figure 3: Diagnostic plots of model3

3. In a manufacturing plant, filters are used to remove pollutants. We are interested in comparing the lifespan of 5 different types of filters. Six filters of each type are tested, and the time to failure in hours is given in the dataset (on the website) `filters` (in `csv` format).

- (a) Use the `read.csv` function to read the data. Then convert the `type` component into a factor.

**Solution:**

```
> filters <- read.csv("data sets/data/filters.csv")
> filters$type <- factor(filters$type)
```

- (b) Using only the `matrix` command, construct a `y` vector and a full rank `X` for this linear model, corresponding to `cont.treatment`

**Solution:**

```
> y <- filters$life
> X.treatment <- matrix(0,30,5)
> X.treatment[,1] <- 1
> # first type is baseline
> for (i in 2:5) { X.treatment [filters$type==i,i] <- 1 }
```

- (c) Fit the models and compare with the `lm` output.

**Solution:**

```
> XtXinv <- solve(t(X.treatment) %*% X.treatment)
> solve(t(X.treatment) %*% X.treatment) %*% t(X.treatment) %*% y
      [,1]
[1,] 249.16667
[2,] -61.66667
[3,] -83.16667
[4,] 108.16667
[5,] 112.16667
```

```

> model <- lm(y~type, data=filters)
> betahat <- model$coefficients
(Intercept)      type2      type3      type4      type5
  249.16667   -61.66667   -83.16667   108.16667   112.16667

```

Both methods give same the same estimates.

- (d) Calculate  $s^2$  using the residuals

**Solution:**

```

> s2 <- sum(lm(y~type, data=filters)$residuals^2)/25
[1] 15304.2

```